

# Class 10: Halloween Mini-Project

Selma Cifric (PID A69042976)

## Table of contents

Importing candy data . . . . .	1
What is your favorite candy? . . . . .	3
Taking a look at pricepercent . . . . .	13
Exploring the correlation structure . . . . .	15
Principal Component Analysis . . . . .	16

## Importing candy data

```
candy_file <- "candy-data.csv"

candy = read.csv("candy-data.csv", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent			
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			
One quarter	0	0	0	0.011	0.511	46.11650			
Air Heads	0	0	0	0.906	0.511	52.34146			
Almond Joy	0	1	0	0.465	0.767	50.34755			

```
library(flextable)
flextable::flextable(head(candy))
```

chocolate	fruity	caramel	peanut	almond	nougat	crispedrice	wafer	hard	bar	pluribus	s
1	0	1	0	0	0	1	0	0	1	0	0
1	0	0	0	0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	1	0	0

```
library(dplyr)
```

The functions dim(), nrow(), table() and sum() may be useful for answering the first 2 questions.

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

There are 85 different candies in this dataset.

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

There are 38 fruity candy types in this dataset.

## **What is your favorite candy?**

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

My favorite candy is Sour Patch Kids.

```
candy["Sour Patch Kids", ]$winpercent
```

```
[1] 59.864
```

Their win percent value is 59.864.

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Winpercent value for Kit Kat is 76.76.

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

Winpercent value for Tootsie Roll Snack Bars is 49.65.

Side-note: the `skimr::skim()` function. There is a useful `skim()` function in the `skimr` package that can help give you a quick overview of a given dataset. Let’s install this package and try it on our candy data.

```
library("skimr")
skim(candy)
```

Table 2: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

**Variable type: numeric**

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

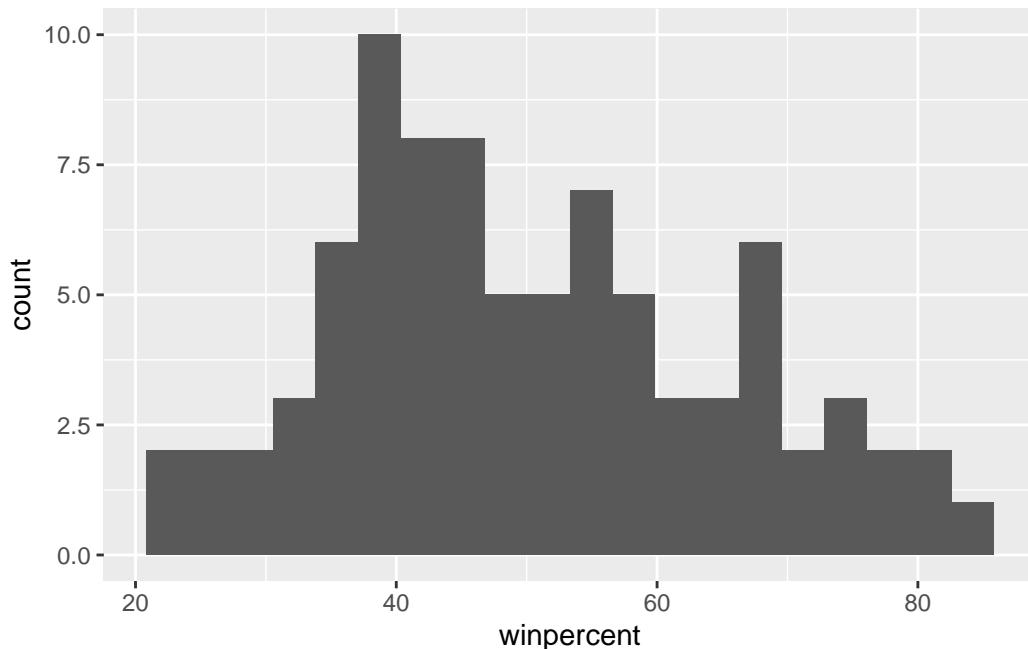
winpercent mean is much higher than the rest of the variables' means, by 50%, and it doesn't use 0-1 scale like others.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

0 or 1 mean yes or no, where 1 means that candy has the chocolate and 0 means no chocolate present in that candy.

Q8. Plot a histogram of winpercent values

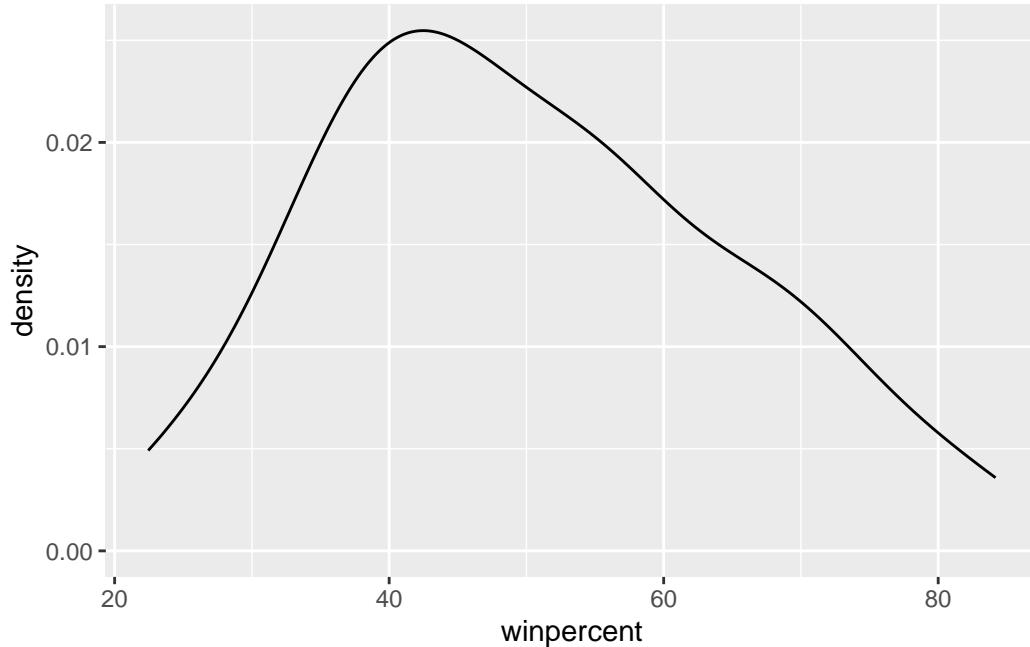
```
library(ggplot2)
ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=20)
```



Q9. Is the distribution of winpercent values symmetrical?

No, the distribution is slightly skewed to the left.

```
ggplot(candy) +
  aes(winpercent) +
  geom_density()
```



Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

```
[1] 50.31676
```

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

From median, we can tell the distribution is under 50% which could be a better measure of centrality than mean. The mean suggests that the center of distribution is right at 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
# 1. Find all chocolate candy in the dataset.
choc inds <- as.logical(candy$chocolate)
choc candy <- candy[choc inds, ]
head(choc candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat		
100 Grand	1	0	1		0	0		
3 Musketeers	1	0	0		0	1		
Almond Joy	1	0	0		1	0		
Baby Ruth	1	0	1		1	1		
Charleston Chew	1	0	0		0	1		
Hershey's Kisses	1	0	0		0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
100 Grand		1	0	1		0	0.732	0.860
3 Musketeers		0	0	1		0	0.604	0.511
Almond Joy		0	0	1		0	0.465	0.767
Baby Ruth		0	0	1		0	0.604	0.767
Charleston Chew		0	0	1		0	0.604	0.511
Hershey's Kisses		0	0	0		1	0.127	0.093
	winpercent							
100 Grand	66.97173							
3 Musketeers	67.60294							
Almond Joy	50.34755							
Baby Ruth	56.91455							
Charleston Chew	38.97504							
Hershey's Kisses	55.37545							

```
# 2. Extract their `winpercent` values.
```

```
choc.win <- choc.candy$winpercent
choc.win
```

```
[1] 66.97173 67.60294 50.34755 56.91455 38.97504 55.37545 62.28448 56.49050
[9] 59.23612 57.21925 76.76860 71.46505 66.57458 55.06407 73.09956 60.80070
[17] 64.35334 47.82975 54.52645 70.73564 66.47068 69.48379 81.86626 84.18029
[25] 73.43499 72.88790 65.71629 34.72200 37.88719 76.67378 59.52925 48.98265
[33] 43.06890 45.73675 49.65350 81.64291 49.52411
```

```
# 3. Find the mean of these values.
```

```
mean(choc.win)
```

```
[1] 60.92153
```

```
# 4-6. Do the same for the fruity candy.
```

```
fruit inds <- as.logical(candy$fruity)
fruit.candy <- candy[fruit inds, ]
head(fruit.candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	
Air Heads	0	1	0		0	0	
Caramel Apple Pops	0	1	1		0	0	
Chewey Lemonhead Fruit Mix	0	1	0		0	0	
Chiclets	0	1	0		0	0	
Dots	0	1	0		0	0	
Dum Dums	0	1	0		0	0	
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent
Air Heads	0	0	0		0	0.906	
Caramel Apple Pops	0	0	0		0	0.604	
Chewey Lemonhead Fruit Mix	0	0	0		1	0.732	
Chiclets	0	0	0		1	0.046	
Dots	0	0	0		1	0.732	
Dum Dums	0	1	0		0	0.732	
	price	percent	win	percent			
Air Heads	0.511	52.34146					
Caramel Apple Pops	0.325	34.51768					
Chewey Lemonhead Fruit Mix	0.511	36.01763					
Chiclets	0.325	24.52499					
Dots	0.511	42.27208					
Dum Dums	0.034	39.46056					

```
fruit.win <- fruit.candy$winpercent
fruit.win
```

```
[1] 52.34146 34.51768 36.01763 24.52499 42.27208 39.46056 43.08892 39.18550
[9] 46.78335 57.11974 51.41243 42.17877 28.12744 41.38956 39.14106 52.91139
[17] 46.41172 55.35405 22.44534 39.44680 41.26551 37.34852 35.29076 42.84914
[25] 63.08514 55.10370 45.99583 59.86400 52.82595 67.03763 34.57899 27.30386
[33] 54.86111 48.98265 47.17323 45.46628 39.01190 44.37552
```

```
mean(fruit.win)
```

```
[1] 44.11974
```

```
# 7. Which mean value is higher?  
mean(choc.win)>mean(fruit.win)
```

```
[1] TRUE
```

Yes, chocolate candy is higher ranked than fruity candy.

Q12. Is this difference statistically significant?

Hint: The chocolate, fruity, nougat etc. columns indicate if a given candy has this feature (i.e. one if it has nougat, zero if it does not etc.). We can turn these into logical (a.k.a. TRUE/FALSE) values with the `as.logical()` function. We can then use this logical vector to access the corresponding candy rows (those with TRUE values). For example to get the `winpercent` values for all nougat containing candy we can use the code: `candywinpercent[as.logical(candy.nougat)]`. In addition the functions `mean()` and `t.test()` should help you answer the last two questions here.

```
t.test(choc.win, fruit.win)
```

Welch Two Sample t-test

```
data: choc.win and fruit.win  
t = 6.2582, df = 68.882, p-value = 2.871e-08  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 11.44563 22.15795  
sample estimates:  
mean of x mean of y  
 60.92153 44.11974
```

Yes, the difference is statistically significant indicated by the p-value of `2.871e-08`.

Q13. What are the five least liked candy types in this set?

Example of using `order()` function.

```
x <- c(10, 2, 5, 1)  
order(x)
```

```
[1] 4 2 3 1
```

Using `head()` and `order()` function we can see the least favorite candy.

```
ord.ind <- order(candy$winpercent)
head(candy[ord.ind, ], 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	
Nik L Nip	0	1	0		0	0	
Boston Baked Beans	0	0	0		1	0	
Chiclets	0	1	0		0	0	
Super Bubble	0	1	0		0	0	
Jawbusters	0	1	0		0	0	
	crispedrice	wafers	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1		0.197	0.976
Boston Baked Beans	0	0	0	1		0.313	0.511
Chiclets	0	0	0	1		0.046	0.325
Super Bubble	0	0	0	0		0.162	0.116
Jawbusters	0	1	0	1		0.093	0.511
	winpercent						
Nik L Nip	22.44534						
Boston Baked Beans	23.41782						
Chiclets	24.52499						
Super Bubble	27.30386						
Jawbusters	28.12744						

These are the least 5 liked candy: Nik L Nip, Boston Baked Bean, Chiclets, Super bubble, and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?

```
tail(candy[ord.ind, ], 5)
```

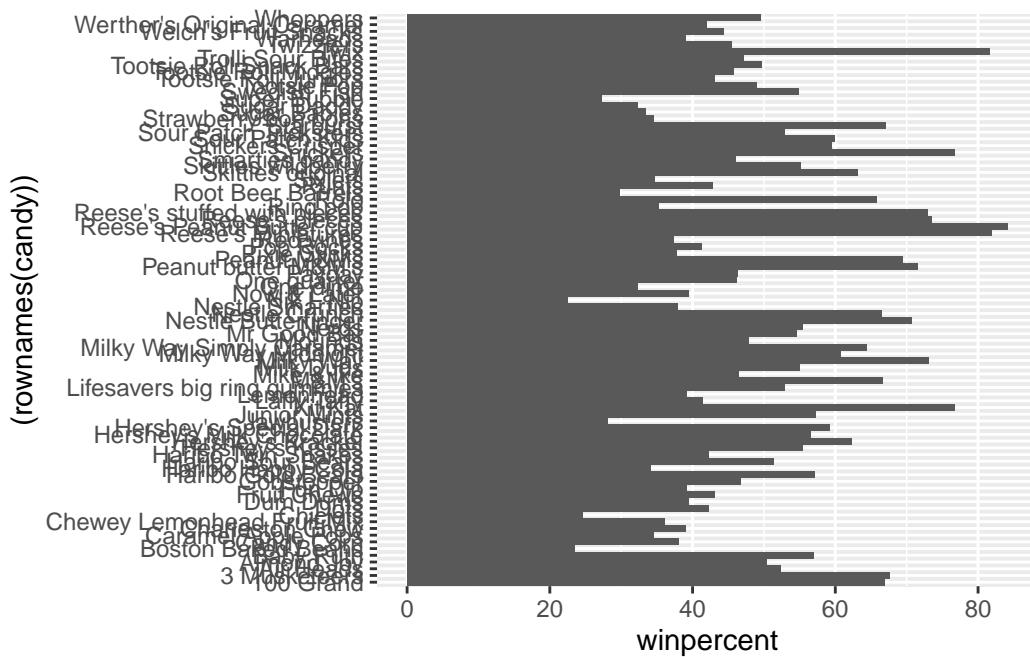
	chocolate	fruity	caramel	peanuty	almondy	nougat	
Snickers	1	0	1		1	1	
Kit Kat	1	0	0		0	0	
Twix	1	0	1		0	0	
Reese's Miniatures	1	0	0		1	0	
Reese's Peanut Butter cup	1	0	0		1	0	
	crispedrice	wafers	hard	bar	pluribus	sugarpercent	
Snickers	0	0	1	0		0.546	
Kit Kat	1	0	1	0		0.313	
Twix	1	0	1	0		0.546	

Reese's Miniatures	0	0	0	0	0.034
Reese's Peanut Butter cup	0	0	0	0	0.720
	price	percent	winpercent		
Snickers	0.651	76.67378			
Kit Kat	0.511	76.76860			
Twix	0.906	81.64291			
Reese's Miniatures	0.279	81.86626			
Reese's Peanut Butter cup	0.651	84.18029			

Top 5 most liked candy are Snickers, Kit Kat, Twin, Reese's Miniatures, and Reese's Peanut Butter Cups.

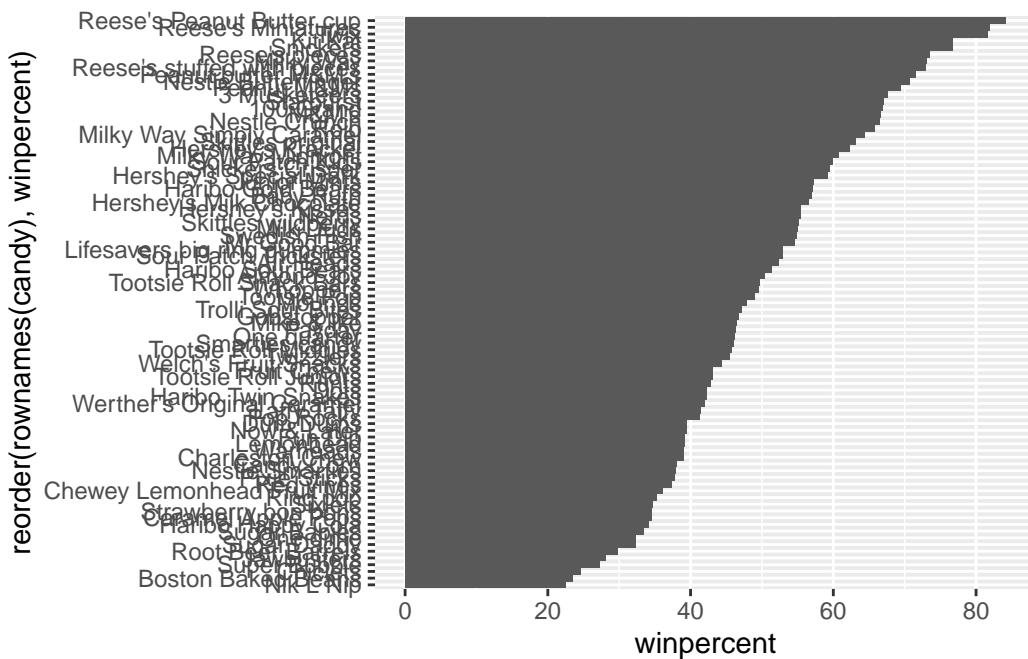
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +
  aes(winpercent, (rownames(candy))) +
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



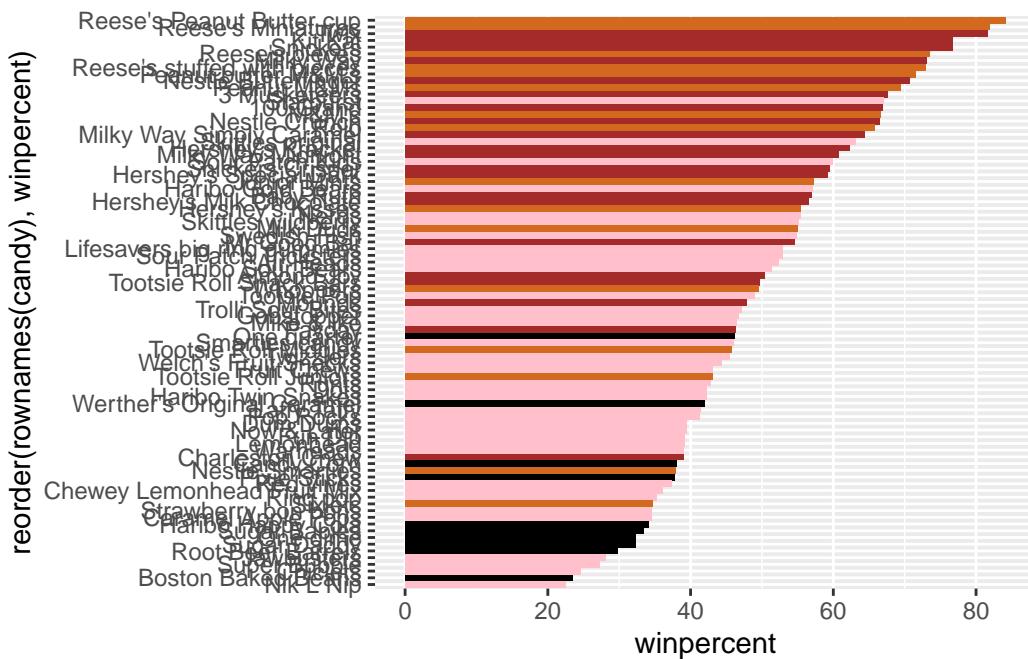
Let's add some color to this plot:

```
my_cols <- rep("black", nrow(candy))
my_cols[candy$chocolate==1] <- "chocolate"
my_cols[candy$bar==1] <- "brown"
my_cols[candy$fruity==1] <- "pink"
my_cols
```

```
[1] "brown"      "brown"       "black"        "black"        "pink"        "brown"
[7] "brown"      "black"       "black"        "pink"        "brown"       "pink"
[13] "pink"       "pink"        "pink"         "pink"        "pink"        "pink"
[19] "pink"       "black"       "pink"         "pink"        "chocolate"   "brown"
[25] "brown"      "brown"       "pink"         "chocolate"   "brown"       "pink"
[31] "pink"       "pink"        "chocolate"    "chocolate"   "pink"        "chocolate"
[37] "brown"      "brown"       "brown"        "brown"       "brown"       "pink"
[43] "brown"      "brown"       "pink"         "pink"        "brown"       "chocolate"
[49] "black"      "pink"        "pink"         "chocolate"   "chocolate"   "chocolate"
[55] "chocolate"  "pink"        "chocolate"    "black"       "pink"        "chocolate"
```

```
[61] "pink"      "pink"      "chocolate" "pink"      "brown"      "brown"
[67] "pink"      "pink"      "pink"      "pink"      "black"      "black"
[73] "pink"      "pink"      "pink"      "chocolate" "chocolate"  "brown"
[79] "pink"      "brown"     "pink"      "pink"      "pink"      "black"
[85] "chocolate"
```

```
ggplot(candy) +
  aes(x=winpercent,
      y=reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols)
```



Q17. What is the worst ranked chocolate candy?

Sixlets are the least favorite chocolate candy.

Q18. What is the best ranked fruity candy?

Starbursts are the best ranked fruity candy (top most pink bar).

## Taking a look at pricepercent

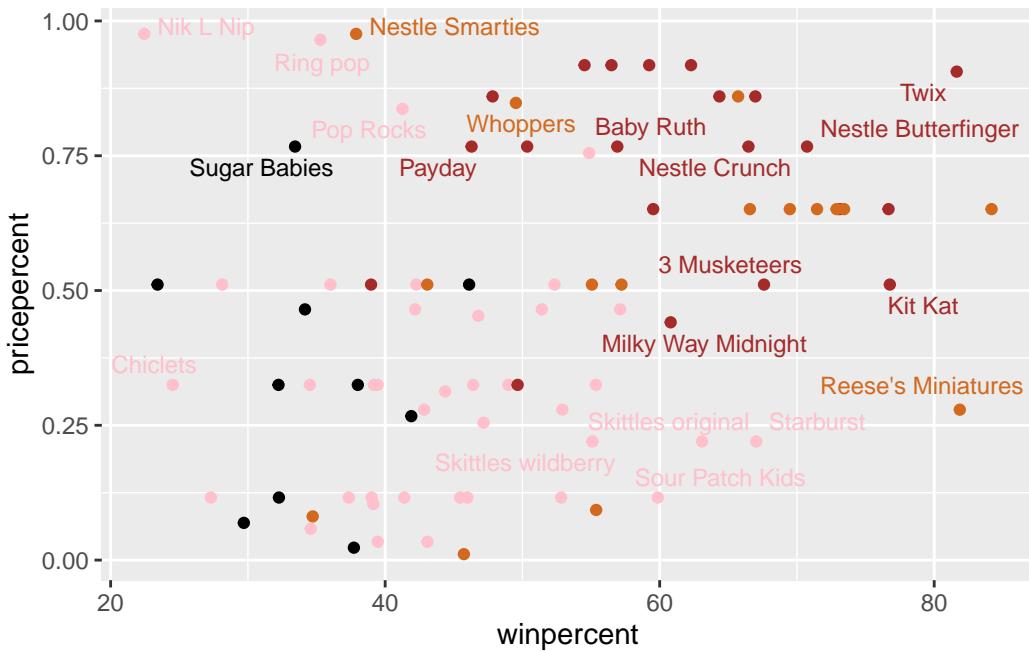
```

library(ggrepel)

ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)

```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

It's chocolate candy Reese's Miniatures.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```

ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )

```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Top 5 most expensive but least popular candy are: Nik L Nip, Nestle Smarties, Ring pop, Hershey's Krackel, and Hershey's Milk Chocolate.

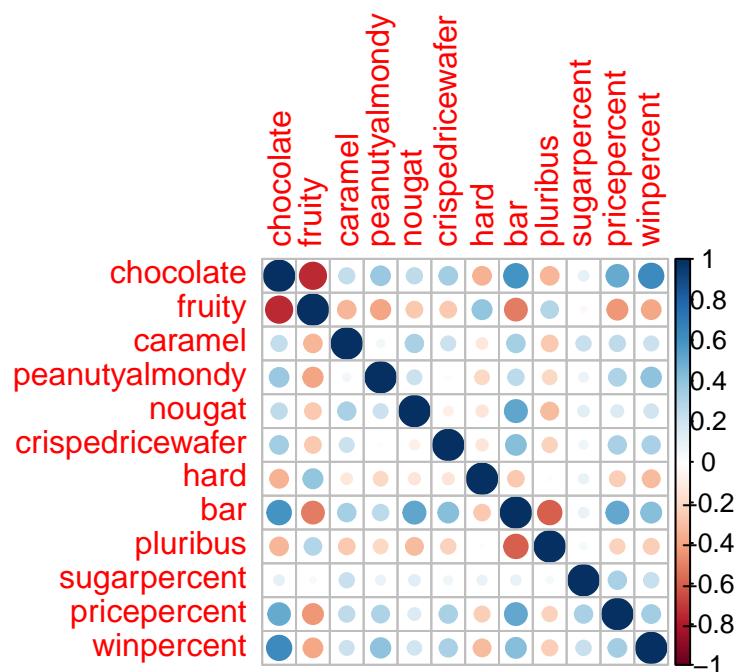
### Exploring the correlation structure

```
cij <- cor(candy)
```

```
library(corrplot)
```

corrplot 0.95 loaded

```
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity candy are anti-correlated.

Q23. Similarly, what two variables are most positively correlated? HINT: Do you like chocolaty fruity candies?

Chocolate and winpercent as well as chocolate and bar. I actually like a nice dark chocolate raspberry mix, which seems to be an unpopular opinion.

## Principal Component Analysis

The main function in base R for this is `prcomp()` and we want to set the `scale=T` here:

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

Importance of components:

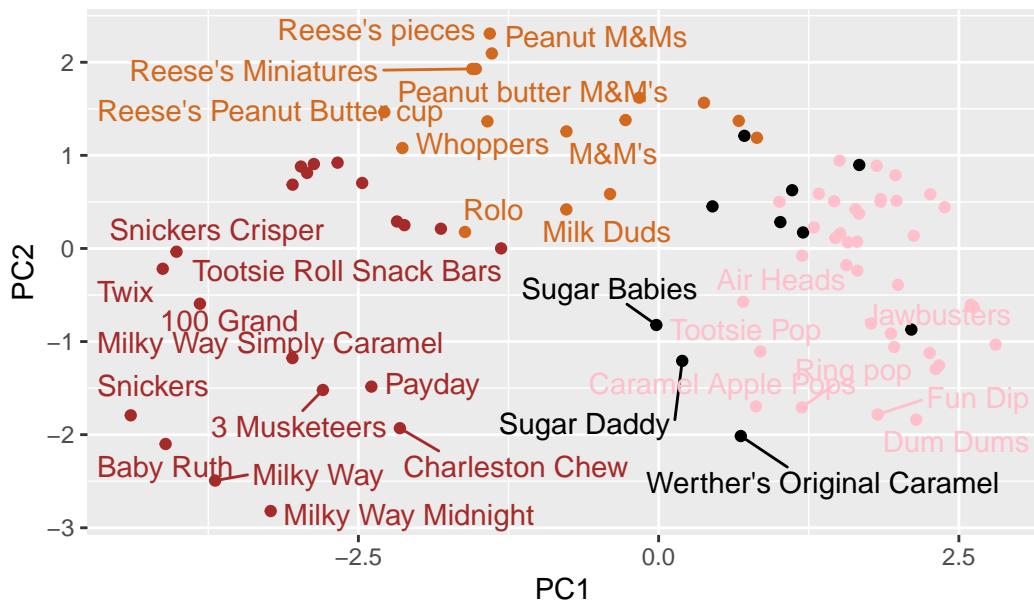
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Let's look at our first main result figure - the "PCA plot" or PC1 vs PC2

```
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols) +
  labs(title="Selma's candy map")
```

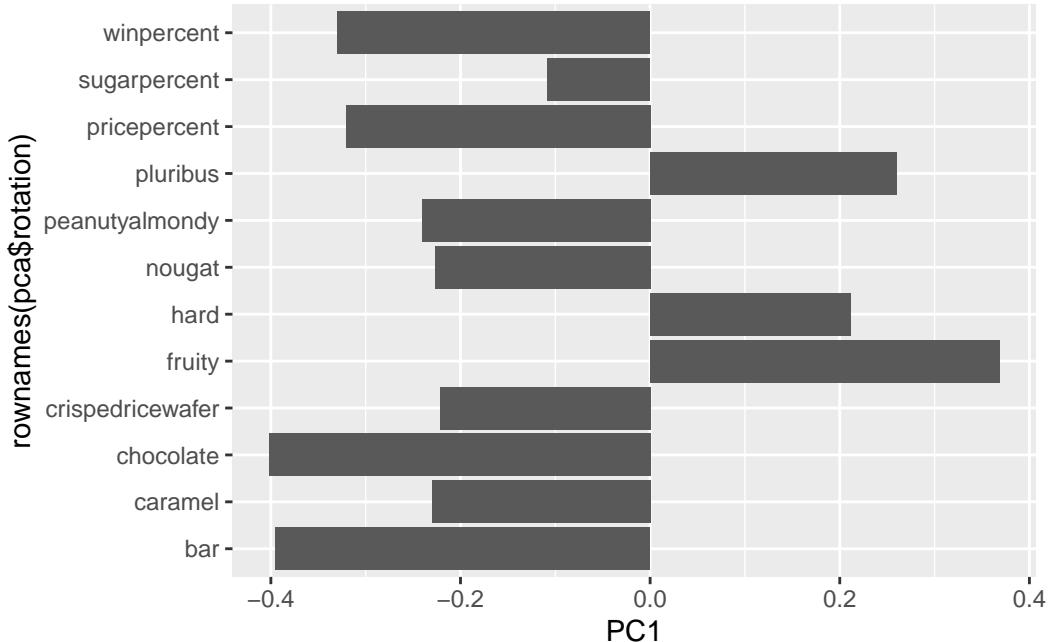
Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Selma's candy map



Don't forget about your variable "loadings" - how the original variable contribute to your new PCs...

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

Fruity, hard and pluribus are picked up strongly by PC1. It makes sense, most fruity candy are hard and come in a pack of multiples.