

# Class 9: Structural Bioinformatics pt. 1

Selma Cifric (PID A69042976)

## Table of contents

The PDB format . . . . .	1
OPTIONAL - extra credit - generating a plot . . . . .	3
Visualizing structure data . . . . .	6
Bio3D package for structural bioinformatics . . . . .	8
Prediction of functional motions . . . . .	13
Play with 3D viewing in R . . . . .	14
<b>Install packages in the R console NOT your Rmd/Quarto file</b>	<b>15</b>

## The PDB format

The main repository for biomolecular structure data is the Protein Data Bank (PDB)  
<https://www.rcsb.org>

Let's have a quick look at the composition of this database:

```
raw.data <- "Data Export Summary.csv"  
protein.df <- read.csv(raw.data)
```

```
protein.df
```

	Molecular.Type	X.ray	EM	NMR	Integrative	Multiple.methods
1	Protein (only)	176,378	20,438	12,709	342	221
2	Protein/Oligosaccharide	10,284	3,396	34	8	11
3	Protein/NA	9,007	5,931	287	24	7
4	Nucleic acid (only)	3,077	200	1,554	2	15
5	Other	174	13	33	3	0
6	Oligosaccharide (only)	11	0	6	0	1

	Neutron	Other	Total
1	83	32	210,203
2	1	0	13,734
3	0	0	15,256
4	3	1	4,852
5	0	0	223
6	0	4	22

We can substitute a comma with a space to change the chr to number.

```
as.numeric(sub(", ", "", protein.df$X.ray) )
```

```
[1] 176378 10284 9007 3077 174 11
```

This is annoying, let's try a different import function from the **readr** package. We install this package into our console.

```
library(readr)
stats <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 9
-- Column specification -----
Delimiter: ","
chr (1): Molecular Type
dbl (4): Integrative, Multiple methods, Neutron, Other
num (4): X-ray, EM, NMR, Total

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
stats
```

```
# A tibble: 6 x 9
`Molecular Type` `X-ray`    EM    NMR Integrative `Multiple methods` Neutron
<chr>            <dbl> <dbl> <dbl>        <dbl>             <dbl> <dbl>
1 Protein (only) 176378 20438 12709        342            221   83
2 Protein/Oligosacch~ 10284 3396   34          8            11    1
3 Protein/NA       9007  5931   287         24            7    0
4 Nucleic acid (only) 3077   200  1554         2            15    3
5 Other             174    13   33          3            0    0
6 Oligosaccharide (o~ 11     0    6           0            1    0
# i 2 more variables: Other <dbl>, Total <dbl>
```

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy.

Percent x-ray

```
n.total <- sum(stats$Total)
n.xray <- sum(stats$`X-ray`)

round((n.xray/n.total)*100, digits = 2)
```

[1] 81.43

81.43% of structures in the PDB are solved by X-Ray.

Percent EM

```
n.em <- sum(stats$EM)
round((n.em/n.total)*100, digits = 2)
```

[1] 12.27

12.27% of structures in the PDB are solved by EM.

Q2: What proportion of structures in the PDB are protein?

```
n.prot <- sum(stats$Total[1])
row.total <- sum(stats$Total)
round((n.prot/row.total)*100, digits = 2)
```

[1] 86.05

86.05% of structures in the PDB are protein.

Q3: Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB?

As of November 1 2025 there are 1150 structures of HIV-1 proteases (no filters applied).


**OPTIONAL - extra credit - generating a plot**

```

library(ggplot2)

ggplot(stats) +
  aes(reorder(`Molecular Type`, - Total), Total, fill = `Molecular Type`) +
  geom_col() +
  labs(x = "Molecular Type", y = "Total", title = "PDB structures") +
  theme(axis.text.x = element_text (angle = 45, hjust = 1))

```



Converting wide to long data frame:

```

library(ggplot2)
library(dplyr)

```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(tidyr)
library(readr)
library(forcats)

# Convert to long format and clean up the data
data_long <- stats %>%
  pivot_longer(
    cols = -c(`Molecular Type`, Total),
    names_to = "experimental_method",
    values_to = "count"
  ) %>%
  mutate(
    count = as.numeric(gsub(", ", "", count))
  ) %>%
  select(-Total)

# Calculate total for each Molecular Type and reorder
data_long <- data_long %>%
  group_by(`Molecular Type`) %>%
  mutate(total = sum(count)) %>%
  ungroup() %>%
  mutate(`Molecular Type` = fct_reorder(`Molecular Type`, total, .desc = TRUE))


colors <- c("X-ray" = "#FF69B4",
           "Other" = "#B19CD9",
           "NMR" = "#00BFFF",
           "Neutron" = "#00CED1",
           "Multiple methods" = "#90EE90",
           "Integrative" = "#DAA520",
           "EM" = "#FF6B6B")

# Create the plot
ggplot(data_long, aes(x = `Molecular Type`, y = count, fill = experimental_method)) +
  geom_col() +
  scale_fill_manual(values = colors) +
  scale_y_continuous(labels = scales::comma_format(),
                     breaks = seq(0, 220000, 50000)) +
  labs(
```

```

title = "PDB Structures by Molecular Type and Experimental Method",
x = NULL,
y = "Number of Structures",
fill = "Experimental Method"
) +
theme_minimal() +
theme(
  axis.text.x = element_text(angle = 45, hjust = 1, size = 11),
  axis.text.y = element_text(size = 11),
  axis.title.y = element_text(size = 12, angle = 90),
  plot.title = element_text(size = 14, face = "bold"),
  legend.title = element_text(size = 12, face = "bold"),
  legend.text = element_text(size = 11),
  panel.grid.major.x = element_blank(),
  panel.grid.minor = element_blank()
)


```



## Visualizing structure data

The Mol\* viewer is embeded in many bioinformatics websites <https://molstar.org/viewer/>.

I can insert any figure or image file using markdown format.



Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?


Hydrogen is a small and light element, not easily picked up by X-ray crystallography which is why it does not show in H<sub>2</sub>O structure. Oxygen, on the other hand, is heavier element and shows up in biological structures on PDB.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have

Yes, its HOH 306 shown as a single red ball in the picture below.

Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic

residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.



## Bio3D package for structural bioinformatics

We can use the bio3d package to read an analyze biomolecular data in R:

```
library(bio3d)  
  
hiv <- read.pdb("1HSG")
```

Note: Accessing on-line PDB file

```
hiv
```

Call: read.pdb(file = "1HSG")

Total Models#: 1  
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

```
Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,
calpha, remark, call
```

Q7: How many amino acid residues are there in this pdb object?

There is 198 amino acid residues/Calpha atoms.

Q8: Name one of the two non-protein residues?

HOH 127 and MK1 1.

Q9: How many protein chains are in this structure?

Two, chains A and B.

```
head(hiv$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40
	segid	elesy	charge										
1	<NA>	N	<NA>										
2	<NA>	C	<NA>										
3	<NA>	C	<NA>										
4	<NA>	O	<NA>										
5	<NA>	C	<NA>										
6	<NA>	C	<NA>										

Let's get the sequence.

```
 pdbseq(hiv)
```

```
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20  
"P" "Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K"  
21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40  
"E" "A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G"  
41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60  
"R" "W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D"  
61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80  
"Q" "I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T"  
81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 1  
"P" "V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F" "P"  
2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21  
"Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K" "E"  
22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41  
"A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G" "R"  
42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61  
"W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D" "Q"  
62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81  
"I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T" "P"  
82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99  
"V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F"
```

Let's trim to chain A and get just it's sequence:

```
chainA <- trim.pdb(hiv, chain="A")  
chainA.seq <- pdbseq(chainA)
```

Let's blast chain A.

```
library(forcats)  
blast <- blast.pdb(chainA.seq)
```

```
Searching ... please wait (updates every 5 seconds) RID = GFX45VRW014
```

```
.....
```

```
Reporting 249 hits
```

```
head(blast$hit.tbl)
```


	queryid	subjectids	identity	alignmentlength	mismatches	gapopens	q.start	q.end	s.start	s.end	evalue	bitscore	positives	mlog.evalue	pdb.id	acc
1	Query_6618617	1W5V_A	100.00		99	0	0	1								
2	Query_6618617	2FDE_A	100.00		99	0	0	1								
3	Query_6618617	1AJV_A	100.00		99	0	0	1								
4	Query_6618617	2R38_A	98.99		99	1	0	1								
5	Query_6618617	2R3T_A	98.99		99	1	0	1								
6	Query_6618617	1HXB_A	98.99		99	1	0	1								

Plot a quick overview of blast results.

```
hits <- plot(blast)
```

```
* Possible cutoff values:    135 110 69 -2
Yielding Nhits:          186 238 244 249

* Chosen cutoff value of:   69
Yielding Nhits:          244
```



```
hits$pdb.id
```

```
[1] "1W5V_A" "2FDE_A" "1AJV_A" "2R38_A" "2R3T_A" "1HXB_A" "1BV9_A" "1AAQ_A"
[9] "1AXA_A" "1HVS_A" "1ZP8_A" "2QHC_A" "1A8G_A" "204L_A" "5COK_A" "1TCX_A"
[17] "2Z54_A" "1D4S_A" "1BV7_A" "1BWA_A" "1A9M_A" "2FLE_A" "1ODY_A" "1GNN_A"
[25] "1GNM_A" "5YRS_B" "1HEF_E" "10DX_A" "4QGI_A" "1BVE_A" "2AZ8_A" "1A30_A"
[33] "6DH6_A" "6DHO_A" "2I4D_A" "600S_A" "1RL8_A" "5YRS_A" "1ZSF_A" "2Q64_A"
[41] "6DH3_A" "2NPH_A" "2Q63_A" "1LZQ_A" "1FB7_A" "1G6L_A" "1HIV_A" "600U_A"
[49] "1HVC_A" "2I4V_A" "2AZ9_A" "600T_A" "2P3B_B" "5KAO_A" "2WLO_A" "6OPT_A"
[57] "1IZI_A" "1MRX_A" "2PYM_A" "2PYN_A" "1DMP_A" "4K4P_A" "1LV1_A" "1AID_A"
[65] "1LV1_A" "1ZBG_A" "3TKG_A" "1HVC_A" "5YOK_A" "1G6L_A" "1FGC_C" "3K4V_A"
[73] "3KT5_A" "3KT5_A" "4QLH_A" "4QLH_A" "2F3K_A" "4Q5M_A" "2AOC_A" "3B80_A"
[81] "3VF5_A" "2AVQ_A" "1DW6_C" "1KZK_A" "2HS1_A" "1K6C_A" "1MTB_A" "4Q1X_A"
[89] "4Q1W_A" "4Q5M_A" "3D1X_A" "2AVM_A" "3PWM_A" "3KT2_A" "3KT2_A" "1SDV_A"
[97] "3JWV_A" "3OY4_A" "1A94_A" "2HS2_A" "4EJ8_A" "2FGU_A" "2AVV_A" "3JW2_A"
[105] "3BVA_A" "1FFF_C" "3S43_B" "2NXD_A" "1FG6_C" "1EBK_C" "4Q1Y_A" "3EL4_A"
[113] "1F7A_A" "1K2B_A" "2FGV_A" "1Z8C_A" "2G69_A" "3EL9_A" "30XV_A" "1BDR_A"
[121] "3N3I_A" "3N3I_A" "30XW_A" "3S43_A" "3EM3_A" "3CYW_A" "5KQX_A" "2B60_A"
[129] "7DOZ_A" "1K2C_A" "1MT7_A" "3EM4_A" "4QJ9_A" "1BDL_A" "3LZS_A" "5T84_A"
[137] "4DQB_A" "7DOZ_A" "4QJ2_A" "3LZV_A" "1SGU_A" "2FXE_A" "1BDQ_A" "3U71_A"
[145] "2R5P_A" "4OBD_A" "7MAS_A" "3IXO_A" "3D3T_A" "5YOJ_A" "3LZU_A" "4NJS_A"
[153] "3EKP_A" "1B6J_A" "3EKQ_A" "2RKF_A" "1C6X_A" "7MAR_A" "4DQF_A" "1RPI_A"
```

```
[161] "3OU1_B" "3PJ6_A" "2P3A_A" "60GQ_A" "30Q7_A" "5KR1_A" "30QD_A" "4RVI_A"
[169] "30QA_A" "1B6K_A" "30UD_B" "6MK9_A" "3S09_A" "1Q9P_A" "6I45_A" "7SEP_A"
[177] "4NJT_A" "3BXR_A" "4YOA_A" "4DQC_A" "2FDD_A" "2RKG_A" "4DQH_A" "2P3C_A"
[185] "4EP2_A" "4EP2_A" "4EQO_A" "4NPT_A" "6OPU_A" "4NPU_A" "3U7S_A" "3HAW_A"
[193] "2AZB_A" "3TTP_A" "3HBO_A" "3GGU_A" "7N6T_A" "60PV_A" "4EQO_A" "60PX_A"
[201] "204N_A" "5T2E_A" "3UCB_A" "3KA2_A" "3FSM_A" "60PW_A" "2AZC_A" "3FSM_A"
[209] "3HLO_A" "2P3D_A" "3T3C_A" "7MYP_A" "6054_X" "60PY_A" "4Z4X_A" "60PZ_A"
[217] "2JE4_A" "1DAZ_C" "7MAP_A" "7MAQ_A" "1K1U_A" "2B7Z_A" "3MWS_A" "1K1T_A"
[225] "8DCH_A" "3I2L_A" "6P9A_A" "2FXD_A" "2J9J_A" "3DCK_A" "2J9J_B" "3NXE_A"
[233] "2040_A" "2040_A" "3NXE_A" "3KA2_A" "3HLO_A" "5B18_A" "1SIP_A" "2SAM_A"
[241] "1AZ5_A" "1SIV_A" "1HII_A" "1IVP_A"
```

## Prediction of functional motions

We can run a Normal Mode Analysis (NMA) to predict large scale motions/flexibility/dynamics of any biomolecule that we can read into R.

Let's look ADK and chain A only.

```
adk <- read.pdb("1ake")
```

```
Note: Accessing on-line PDB file
PDB has ALT records, taking A only, rm.alt=TRUE
```

```
adk_A <- trim.pdb(adk, chain="A")
adk_A
```

```
Call: trim.pdb(pdb = adk, chain = "A")

Total Models#: 1
Total Atoms#: 1954, XYZs#: 5862 Chains#: 1 (values: A)

Protein Atoms#: 1656 (residues/Calpha atoms#: 214)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 298 (residues: 242)
Non-protein/nucleic resid values: [ AP5 (1), HOH (241) ]

Protein sequence:
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLVT
```


```
DELVIALVKERIAQEDCRNGFLLDGFPR TIPQADAMKEAGINV D YVLEFDVPDELIVDRI  
VGRRVHAPSGRVYHVKF NPPKVEGKDDVTGEELTRKDDQEETVRKRLVEYHQMTAPLIG  
YY SKEAEAGNTKYAKVDGTPVAEV RADLEKILG
```

```
+ attr: atom, helix, sheet, seqres, xyz,  
       calpha, call
```

```
m <- nma(adk_A)
```

```
Building Hessian...           Done in 0.02 seconds.  
Diagonalizing Hessian...     Done in 0.446 seconds.
```

```
plot(m)
```



Let's write out a “trajectory” of predicted motion.

```
mktrj(m, file="adk_nma.pdb")
```

## Play with 3D viewing in R

We can use the new **bio3Dview** package, which is not yet on CRAN, to render interactive 3D views in R and HTML quarto output reports. To install from GitHub we can use the **pak** package. We install it in the console.

## **Install packages in the R console NOT your Rmd/Quarto file**

```
#install.packages("bio3d") #install.packages("devtools") #install.packages("BiocManager")
#BiocManager::install("msa") #devtools::install_bitbucket("Grantlab/bio3d-view")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

msa

Q11. Which of the above packages is not found on BioConductor or CRAN?:

bio3d-view

Q12. True or False? Functions from the devtools package can be used to install packages from GitHub and BitBucket?

True